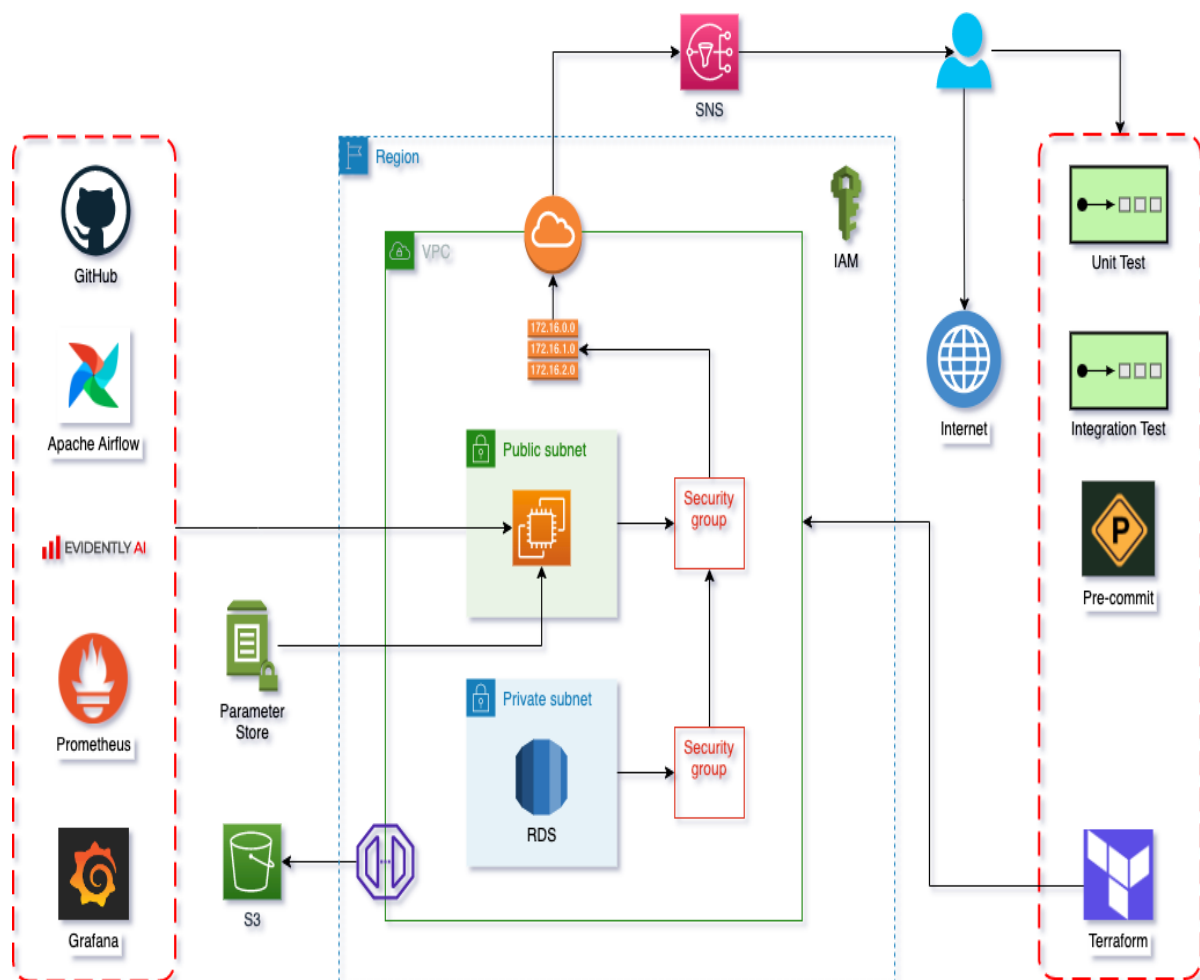


Waiter Tips Prediction Run Book

In this **Waiter Tips Prediction** MLOps project, we predict waiter tips with **XGBoost** based on features named **total bill, gender, smoker, day, time and size**. We first deploy resources on **AWS** with **Terraform**. Once we have all the data, the model, and scripts on the local machine and have completed the tests, we commit the project directory onto **GitHub**, which, through **Actions**, installs the files on Ubuntu virtual machine.

Later, we initiate the **MLFlow** and **Airflow** servers to track experiments and manage the workflow. **Airflow** automatically runs every month, checking if any new data exists in the **S3** bucket and retrieving it from there to train a new model. The application saves the latest model in the **S3** bucket to use in the production environment. We can watch all metrics by **Evidently** on the localhost and by **Prometheus** on the **Grafana** dashboard. Every time we train the model and check for any data and concept drift, the user also receives an email notification.



Users can access the application through a web interface on **Flask**.

We run the tests on the local machine by implementing **pytest, black, isort, localstack, pre-commit, and pylint** before committing to **GitHub**.

Installation

1. Create an AWS account and get a programmatic access.
2. Deploy AWS resources with *Terraform*. Run in the terraform folder:

```
terraform init
```

```
terraform plan
```

```
terraform apply -auto-approve
```

3. Commit the project to *GitHub*:

```
git add .
```

```
git commit -m 'initial commit'
```

```
git push origin main
```

If Actions are not in place, you can use **wget** command.

4. Install **brew**, **Prometheus**, start **Prometheus** and **Grafana** servers, and create **PostgreSQL** databases on **RDS** running the following command in the project folder:

```
start.sh
```

Run

1. Start the **MLFlow** server:

```
mlflow server -h 0.0.0.0 -p 5000 --backend-store-uri
```

```
postgresql://DB_USER:DB_PASSWORD@DB_ENDPOINT:5432/DB_NAME --default-artifact-root s3://s3b-tip-predictor/mlflow/
```

2. Start the Evidently server:

```
mlflow server -h 0.0.0.0 -p 5500 --backend-store-uri
```

```
postgresql://DB_USER:DB_PASSWORD@DB_ENDPOINT:5432/DB_NAME --default-artifact-root s3://s3b-tip-predictor/evidently/
```

You need to choose the database user and password, and get the RDS database endpoint.

3. Open the **airflow.cfg** and set executor as **LocalExecutor**, and **sql_alchemy_conn** as

```
sql_alchemy_conn = postgresql+psycopg2://<user>:<pass>@<host>:5432/<db>
```

4. Initialize **Airflow** database in the project folder (**/app/Waiter-Tips-Prediction**).

```
airflow db init
```

5. Create a user:

```
airflow users create --username <username> --password <password> --  
firstname <firstname> --lastname <lastname> --role Admin --email <email>
```

6. Start the ***Airflow*** server:

airflow webserver -p 8080 -D

7. Start the ***Airflow*** scheduler:

airflow scheduler -D

8. Do port forwarding on VSC for the following ports to access the Web UIs on the local machine:

- **3000: Grafana**
- **3500: Flask**
- **3600: Flask**
- **5000: MLflow**
- **5500: Evidently**
- **8080: Airflow web server**
- **8793: Airflow scheduler**
- **9090: Prometheus**
- **9091: Prometheus**

The App (**prediction**) runs on port **3500**, **Evidently** reports on **3600**, the app's **Prometheus** metrics on **9091**.

To start **MLFlow** and **Airflow** servers readily, you can use Makefile provided that MLflow server endpoints are updated:

make mlflow_5000

make mlflow_5500

make airflow_web

make airflow_scheduler